

# **Master-Praktikum: IoT**

(IN2106, IN4224)

- IoT Core Team -

Final Presentation  
30.07.2018

## Topic

---

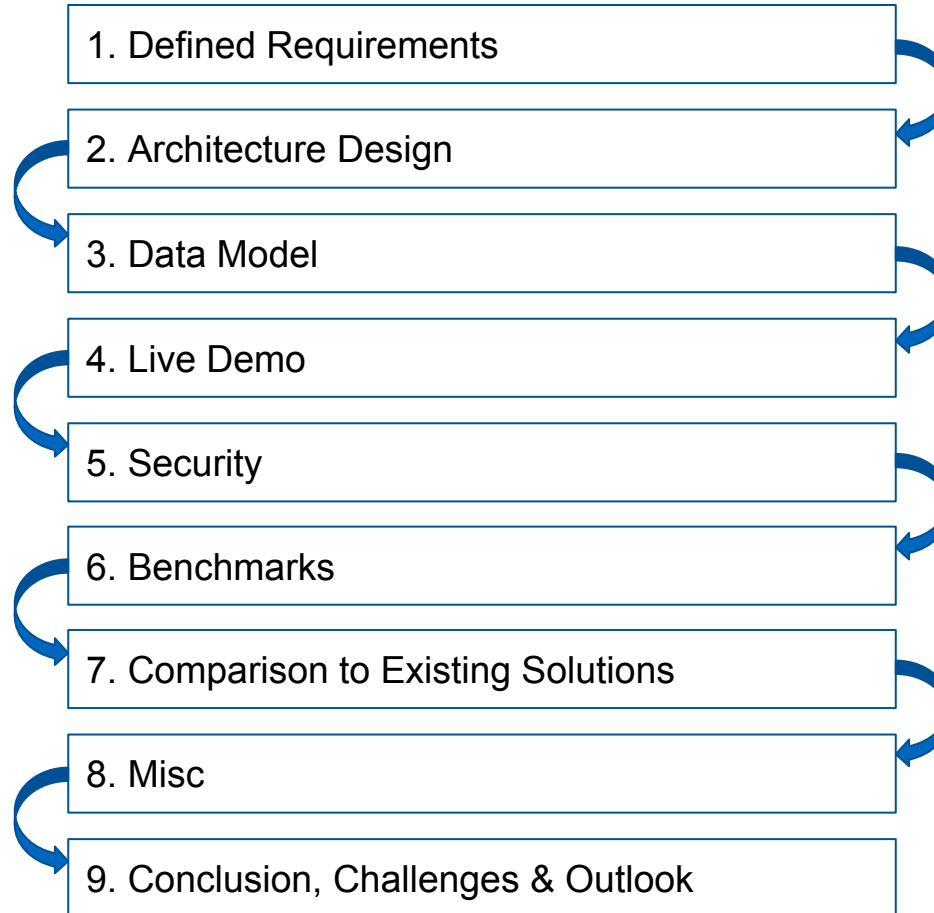
Presenting the design and capabilities of the reinvented IoT platform

## Presenters

---


Peeranut Chindanonda  
Helge Dickel  
Christoph Gebendorfer  
Bahareh Hosseini  
Hans Kirchner

# Agenda



## Defined Requirements

---




## Defined Requirements

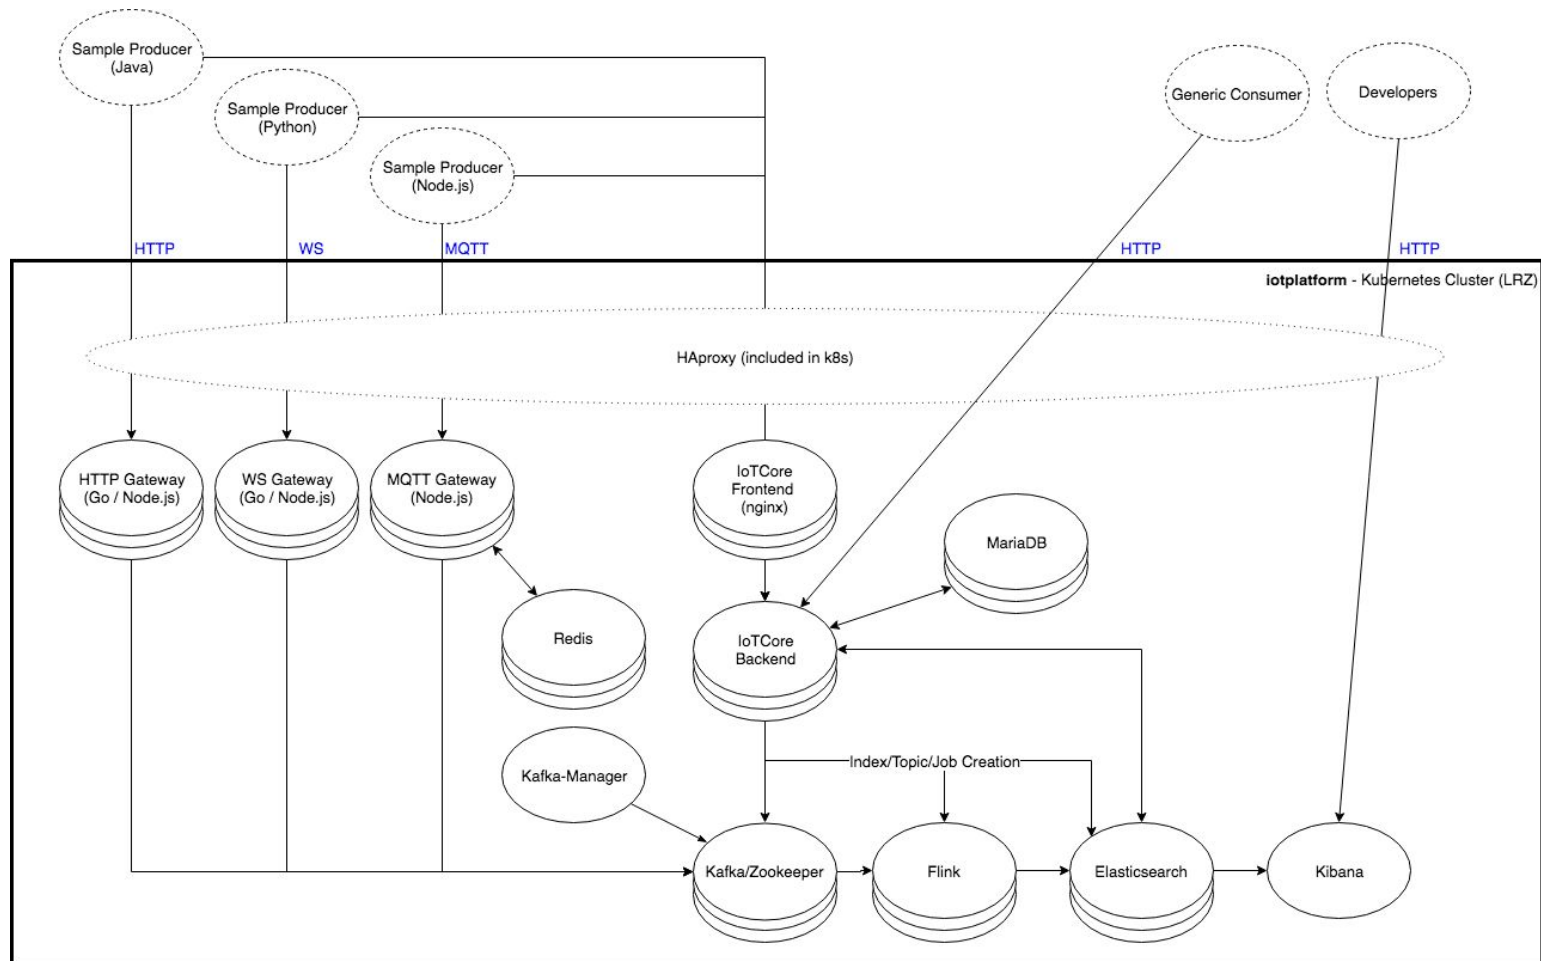
	Requirement	Previous Status
R1	Secure communication & transmission of data	n.a.
R2	Storage	Persisted inside VM
R3	Data Provisioning to Consumers	REST-API
R4	Tested Processing of large amounts of data (>100k msg/sec) backed by scalability and load balancing on all tiers - ingestion to extraction	n.a.
R5	Load throughput testing	n.a.
R6	Ensure platform-independence, platform should be deployable on commonly used OS	n.a.
R7	Ensure that platform can be deployed with beginner knowledge	~ long, detailed guide available
R8	Guide for Deployment & Usage	~ long, detailed guide available
R9	Multitenancy	n.a.
R10	Support Ingestion via MQTT	n.a.

# Architecture Design

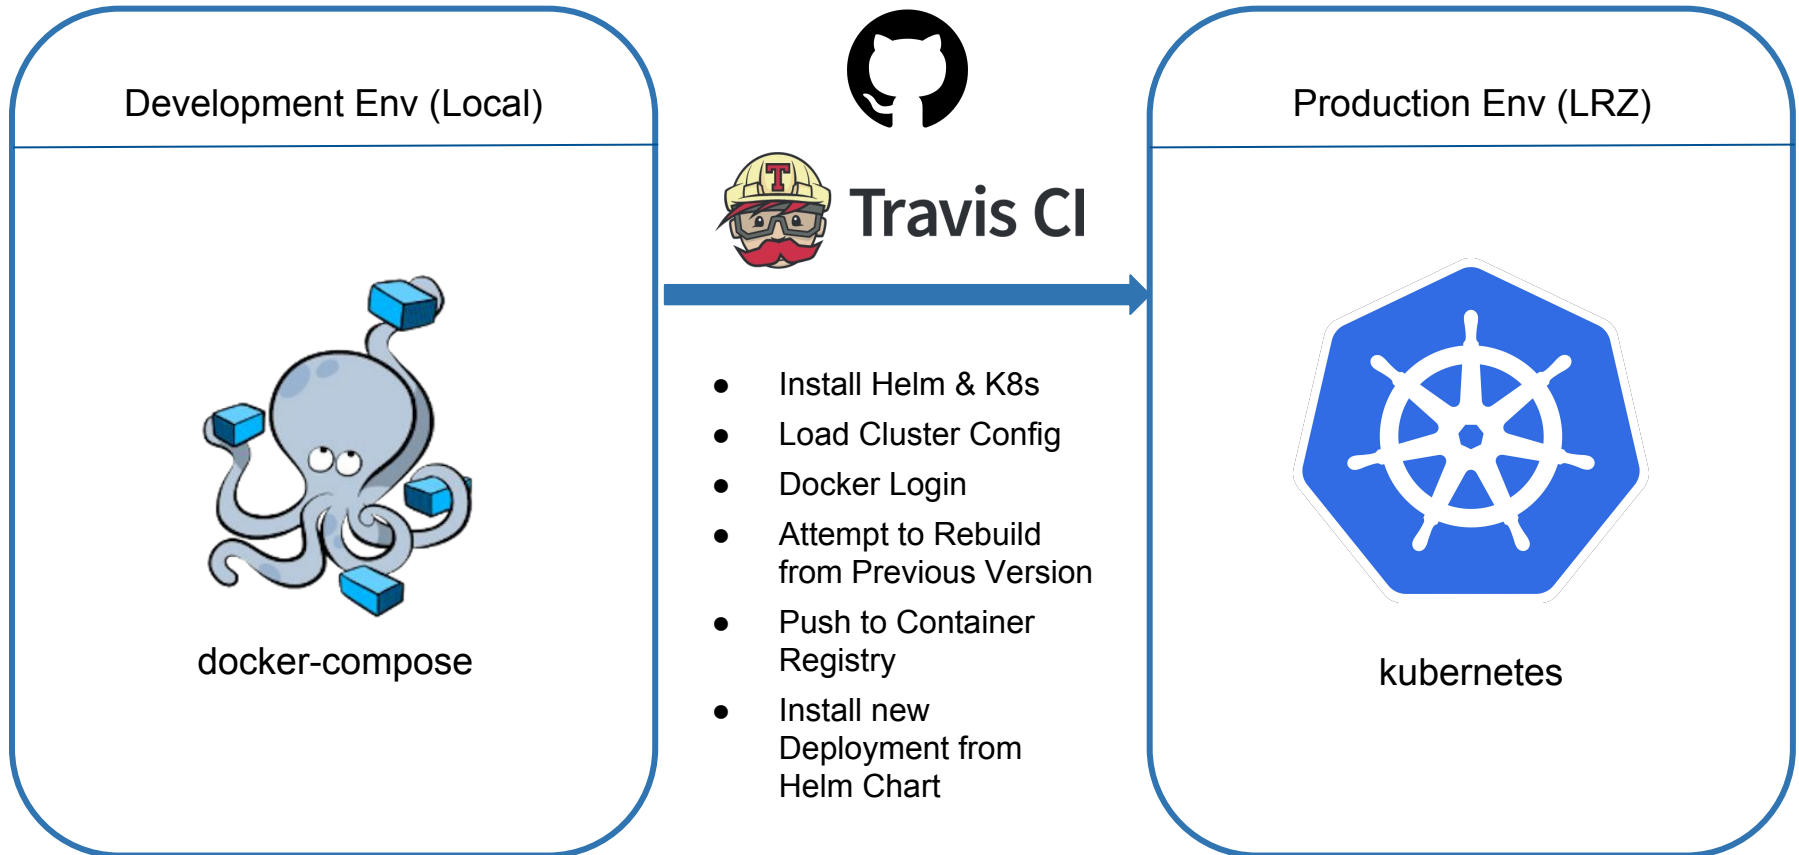
---



# Architecture Design - Pipeline




## Architecture Design - Deployment






# Data Model

---

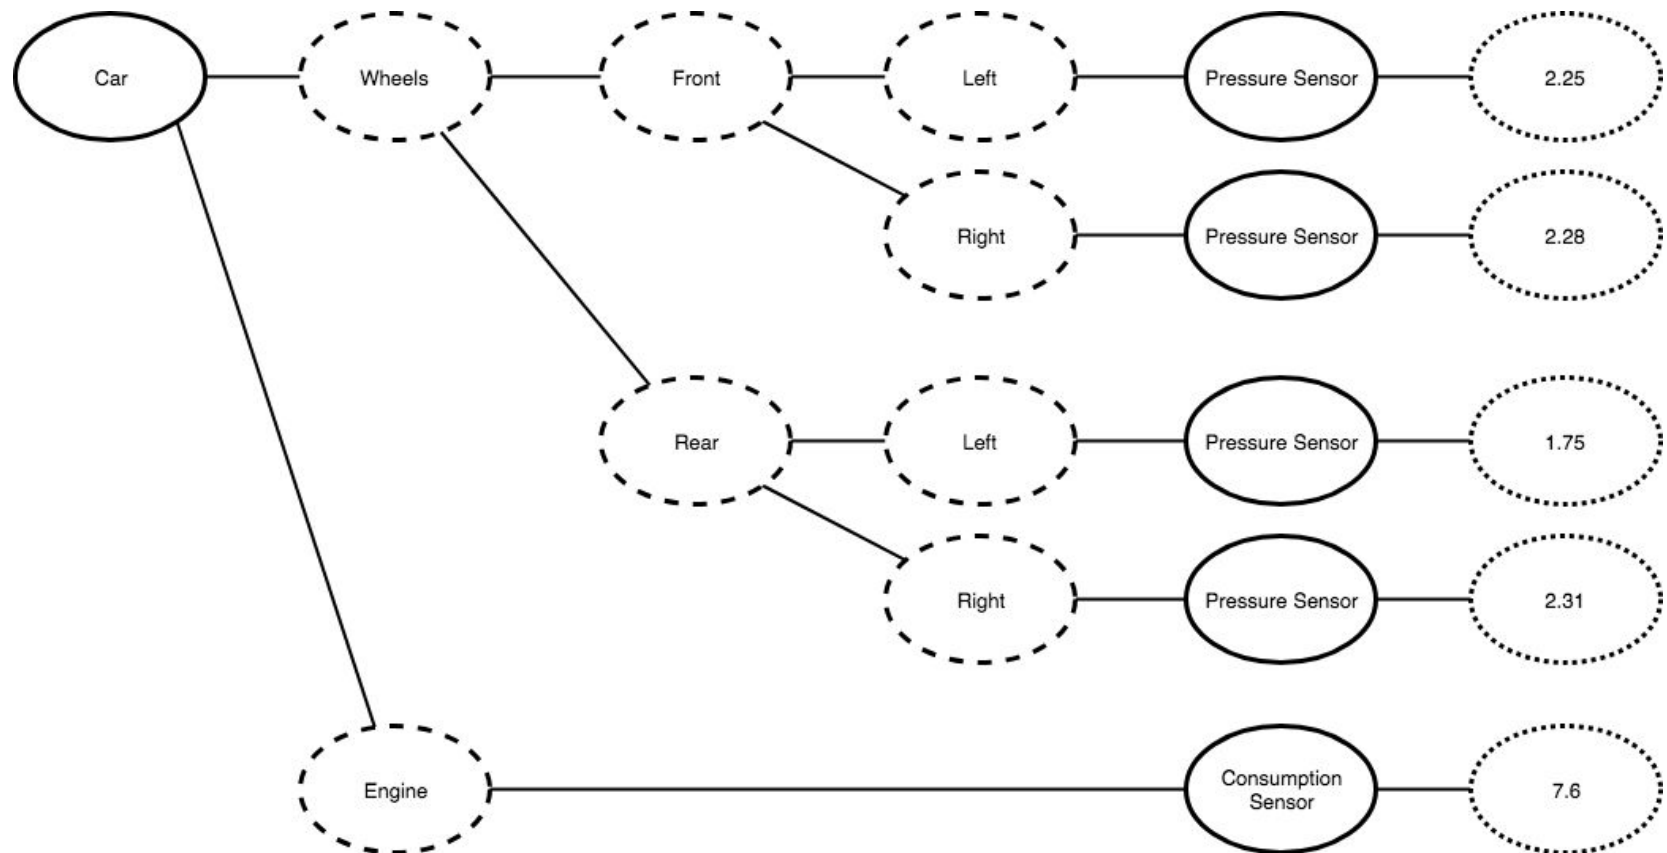


## Data Model - Device-/Sensor Model

- Our Intent
  - Support of generic devices
  - Enable (optional) multi level nesting of sensors
- Device
  - Entity that may contain one or many sensors and represents a single physical logical unit
- Sensor
  - Entity that belongs to a device
  - Represents a single physical measuring point, producing a time series of data




## Data Model - Example




## Live Demo

---



# Security

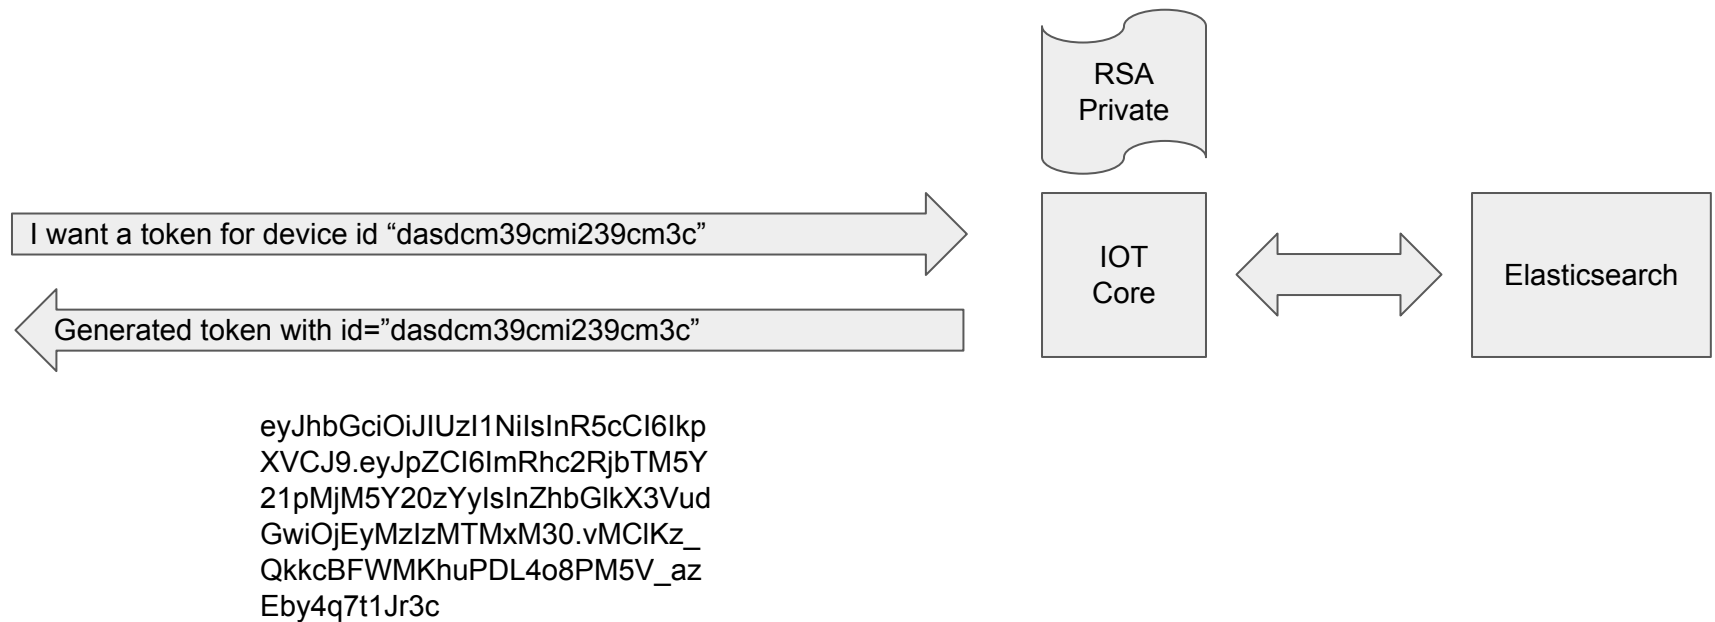
---



# Security

Security via JWT (JSON Web Token, IETF RFC7519 standard) [<https://tools.ietf.org/html/rfc7519>]

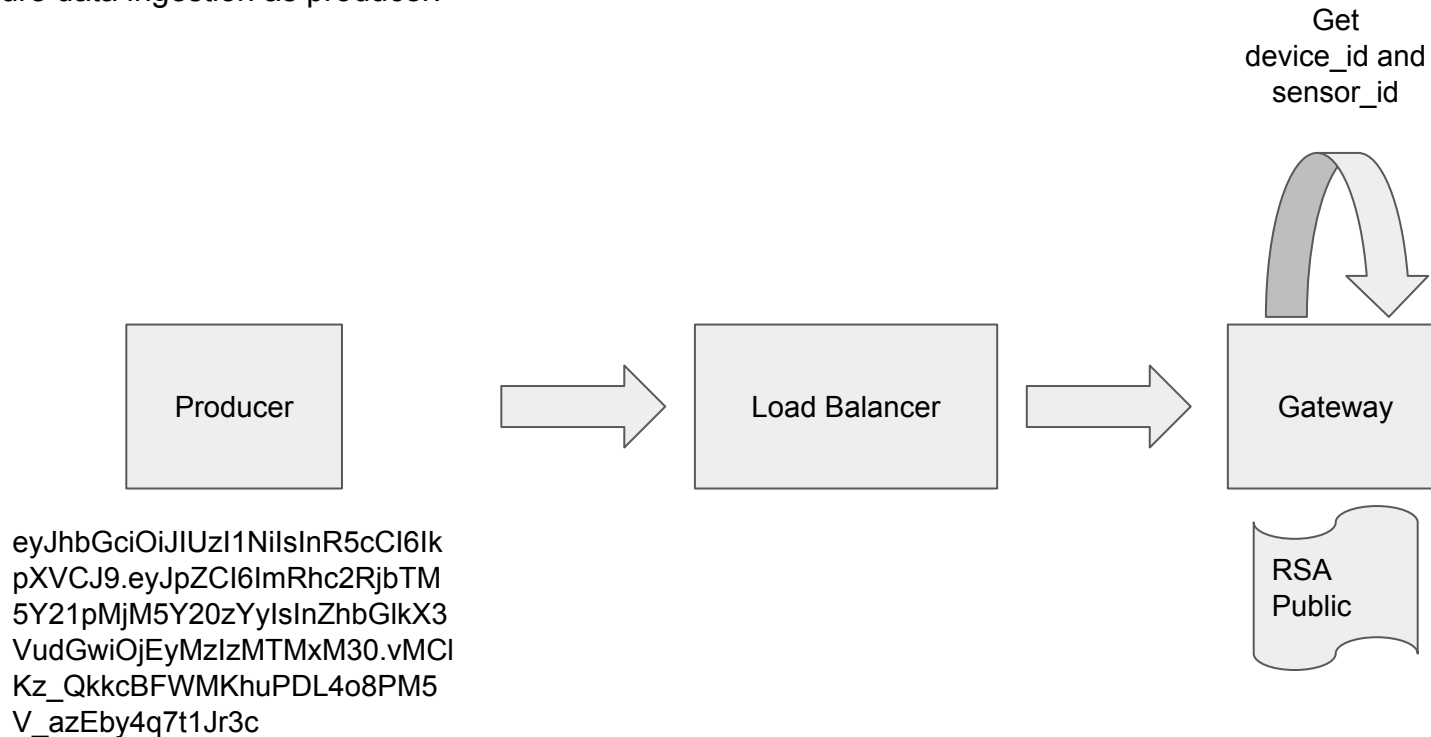
Retrieving JWT Tokens:



# Security


Security via JWT (JSON Web Token, IETF RFC7519 standard) [\[https://tools.ietf.org/html/rfc7519\]](https://tools.ietf.org/html/rfc7519)

Secure data ingestion as producer:



# Benchmarks

---






## Benchmarks - HTTP - w/ Auth

Performance on Hetzner, cx41: 4 vCPUs, 16Gb of RAM

Parameters:


- 100 devices
- 200000 msg
- Max-Throughput



Statistics:

- Mean: 13.236 ms
- Max: 80 ms
- Req/s: 7554.87

Completed Requests



Parameters:


- 100 devices
- 200000 msg
- Max-Throughput



Statistics:

- Mean: 44.942 ms
- Max: 136 ms
- Req/s: 2205.08

Completed Requests




## Benchmarks - HTTP - w/ Auth

Performance on Hetzner, cx41: 4 vCPUs, 16Gb of RAM

Parameters:


- 1000 devices
- 100000 msg
- Max-Throughput



Statistics:

- Mean: 131.057 ms
- Max: 3321 ms
- Req/s: 7630.28

Completed Requests



Parameters:

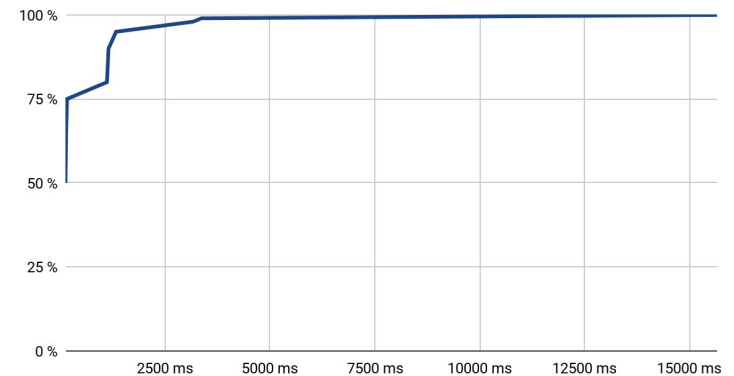
- 1000 devices
- 100000 msg
- Max-Throughput



Statistics:

- Mean: 452.459 ms
- Max: 15662 ms
- Req/s: 2210.14

Completed Requests




# Benchmarks - WebSockets - w/ Auth

Performance on Hetzner, cx41: 4 vCPUs, 16Gb of RAM

Parameters:

- 1000 devices
- 1000 msg/device
- Connections pre-established



Statistics:

- Mean: 0.542 ms
- Max: 923 ms
- Req/s: 800319.53

Parameters:

- 1000 devices
- 1000 msg/device
- Connections pre-established




Statistics:

- Mean: 0.549 ms
- Max: 726 ms
- Req/s: 783266.05


# Benchmarks - WebSockets - Golang, w/ Auth

Performance on Hetzner, cx41: 4 vCPUs, 16Gb of RAM




Parameters:

- n devices
- 1000 msg/device
- Connections pre-established



## Comparison to Existing Solutions

---



# Comparison to Existing Solutions -

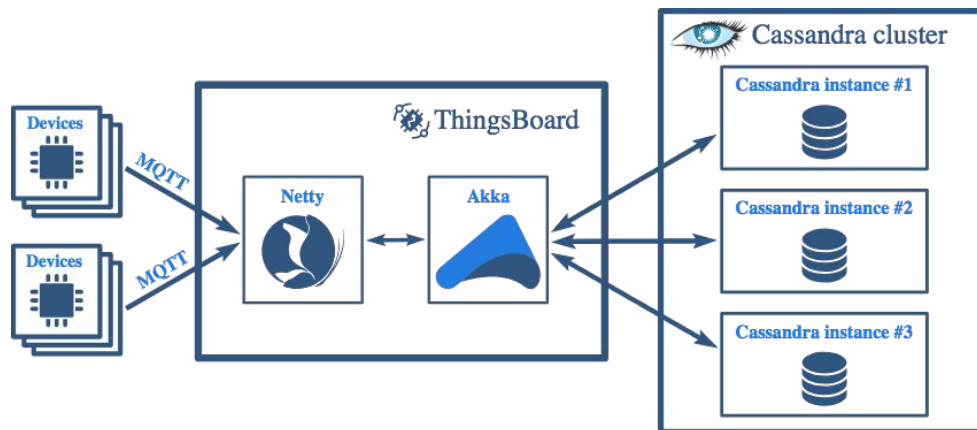
FEATURE	AVOCADO ARCHIPELAGO	BANANA BEACH (Early Access Only)
<b>Architecture</b>	Monolithic	Microservices
<b>Connectivity protocol</b>	Proprietary	Open, standards-based
<b>Gateway connectivity model</b>	One connection per device	Single, multiplexed connection
<b>Communication security</b>	RSA+AES	(D)TLS
<b>Device credential management</b>	No	Yes
<b>Device metadata</b>	Structured	Structured or unstructured
<b>Data collection</b>	Single data type, structured only	Unlimited types, isolated flows, structured or unstructured
<b>Configuration management</b>	Structured only	Structured or unstructured
<b>Data processing and analytics</b>	3-rd party integrations	Built-in or 3-rd party integrations
<b>Data visualization</b>	3-rd party integrations	Built-in customizable dashboards or 3-rd party integrations
<b>Device notifications</b>	Yes	No, superseded by commands
<b>Command execution</b>	No	Yes
<b>Over-the-air updates</b>	No	Yes
<b>Technology stack</b>	Mainly Java	Polylingual
<b>Scalability, elasticity, self-healing</b>	Manual	Automated container orchestration
<b>Server configuration</b>	Non-portable, stored in DB	Portable declarative blueprint

[<https://www.kaaproject.org/whats-new/>]

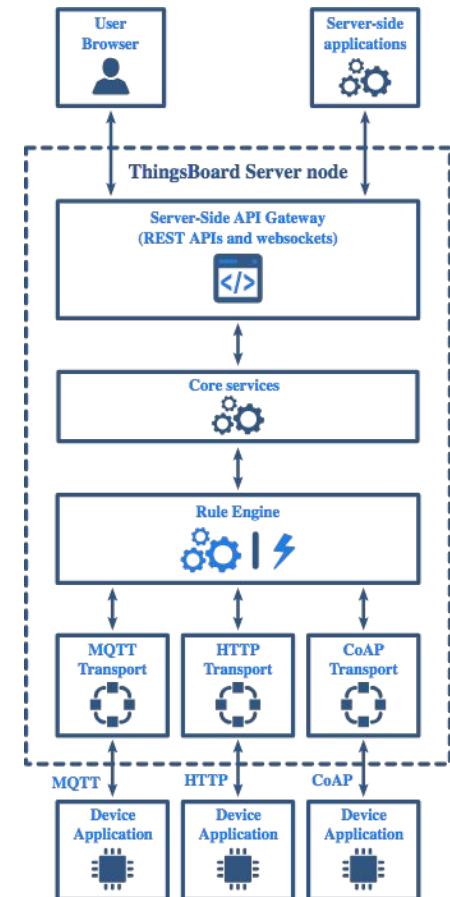
## Comparison to Existing Solutions - ThingsBoard

### Core Services:

- Device and credentials
- Rule chains and rule nodes
- Tenants and customers
- Widgets and dashboards
- Alarms and events



[<https://thingsboard.io/docs/reference/performance/>]



[<https://thingsboard.io/docs/reference/architecture/>]

## Comparison to Existing Solutions - ThingsBoard


AWS, c4.2xlarge: 4 vCPUs, 7.5Gb of RAM

- MQTT
- 10000 devices
- 1 msg/sec/device
- total load: 10000 msg/sec

Statistics:

- Mean: 217 ms
- Max: 10887 ms
- Req/s: 6818.182

Completed Requests




AWS, c4.2xlarge: 8 vCPUs, 15Gb of RAM

- MQTT
- 10000 devices
- 1 msg/sec/device
- total load: 10000 msg/sec

Statistics:

- Mean: 38 ms
- Max: 3270 ms
- Req/s: 8823.529

Completed Requests




[<https://thingsboard.io/docs/reference/performance/>]



## Misc

---




# Misc

adsf

## Conclusion, Challenges & Outlook

---




## Conclusion

	Requirement	Current Status
R1	Secure communication & transmission of data	✓ Messaging & device administration secured via JWT
R2	Storage	✓ Persisted in ES container within cluster
R3	Data provisioning to consumers	✓ Secure REST-API to ES
R4	Processing large amounts of data, scalability in all tiers - ingestion to extraction	✓ Large amounts of data ~ Autoscaling capabilities via K8s
R5	Load throughput testing	✓ Gateway capabilities tested + language comparisons
R6	Ensure platform-independence	✓ Guaranteed thanks to dockerized application design
R7	Ensure accessibility with beginner knowledge	✓ As easy as “docker-compose up”
R8	Guide for deployment & usage	✓ Detailed Github Readme + docs
R9	Multitenancy	✓ Private tokens for devices
R10	Support ingestion via MQTT	✓ MQTT, HTTP & WS gateways available

# Challenges

- Team members had to evolve an understanding of the considerable technology stack
- Ambitious scope
- Opinion: previous architecture design not suited for the requirements given to us
  - Almost inevitable to redesign and reimplement, adding considerable workload
- Cooperation with HAL team
  - Difficult since they have to rely on running architecture, which is hard if it is being reworked
  - Integration now possible
- JWT Authorization - Node.js library converts HTTP headers to lower case
  - Either: loop over raw headers (sacking performance), change header (abusing the standard)



# Outlook

- Activate true persistence, surviving rolling deployments (only Flink missing)
- Finish up on autoscaling (Kafka missing)
- Improve frontend UX
- Security testing
- Actuator expansion (e.g. connected to Flink)
- Provide more default Flink jobs for analytics
- MQTT performance testing



